

Report: Software Engineer Intern Assignment - Zeotap

<https://github.com/ShivamBajpai04/wrangler>

1. Introduction

This report documents the implementation of new features in the CDAP Wrangler library, as per the assignment requirements provided by Zeotap. The core objective was to enhance Wrangler's capabilities by adding native support for parsing and utilizing byte size and time duration units within recipes. This was achieved through the development of `ByteSize` and `TimeDuration` classes, and the introduction of the `AggregateStats` directive to aggregate byte sizes and time durations.

2. Implementation Details

- **2.1 Grammar Modification:**
 - The `Directives.g4` file was modified to incorporate new lexer rules for `BYTE_SIZE` and `TIME_DURATION`.
 - Helper fragments `BYTE_UNIT` and `TIME_UNIT` were added to facilitate the definition of these rules.
 - Parser rules, including `value`, `byteSizeArg`, and `timeDurationArg`, were adjusted to appropriately accept the new `BYTE_SIZE` and `TIME_DURATION` tokens.
 - Following these modifications, the ANTLR Java parser/lexer code was regenerated using the `mvn compile` command.
- **2.2 API Updates:**
 - New Java classes `ByteSize.java` and `TimeDuration.java` were created, extending the `Token` class from the `wrangler-api` module.
 - These classes are designed to parse token strings within their constructors (e.g., "10KB", "150ms").
 - Methods were implemented to retrieve the parsed values in their canonical units; for example, `getBytes()` in the `ByteSize` class.
 - `BYTE_SIZE` and `TIME_DURATION` were added to the Token Types.
 - The usage and token definitions were updated to enable the specification of these new token types as valid directive arguments.
- **2.3 Core Parser Updates:**
 - The `wrangler-core` module was updated by adding visit methods for the parser rules created or modified in the grammar modification step. This includes methods such as `visitByteSizeArg`, `visitTimeDurationArg`, or modifications to `visitValue`.
 - The `ctx.getText()` method was used to extract the text from the parse tree context.
 - Instances of the created tokens were added to the `TokenGroup`.
- **2.4 New Directive Implementation:**
 - A new directive class, `AggregateStats`, was created, implementing the `Directive` interface.
 - The `define()` method of the directive accepts at least four arguments:

- The column name for source byte sizes.
 - The column name for source time durations.
 - The target column name for the total size.
 - The target column name for the total or average time.
- Optional arguments can be used to specify output units (e.g., 'MB', 'GB', 'seconds', 'minutes') or the aggregation type (total, average).
- During initialization, the source and target column names provided in the arguments are stored.
- The directive operates as an aggregate, using a store to accumulate totals.
- In the execution phase:
 - Byte size values are read from the specified source size column.
 - Time duration values are read from the specified source time column.
 - These values are converted to their canonical units (e.g., bytes and nanoseconds) and added to running totals stored in the store.
- The finalization process involves:
 - Retrieving the final totals from the store.
 - Performing unit conversions as required by the arguments.
 - Returning a single new Row containing the target columns with the calculated aggregate values (e.g., `total_size_mb`, `total_time_sec`).

3. Testing

- **3.1 Unit Tests for ByteSize and TimeDuration:**
 - Unit tests were added for the `ByteSize` and `TimeDuration` classes.
 - These tests verify the correct parsing and canonical value retrieval for various inputs, such as "10kb", "1.5MB", "5ms", and "2.1s".
- **3.2 Parser Tests:**
 - Parser tests were implemented (e.g., in `GrammarBasedParserTest.java` or `RecipeCompilerTest.java`).
 - These tests ensure that recipes using the new syntax are parsed correctly and that invalid syntax is rejected.
- **3.3 Unit Tests for AggregateStats Directive:**
 - Comprehensive unit tests were developed for the `AggregateStats` directive.
 - The test case specification includes:
 - Input Data: `Row` objects were created to simulate sample log or transaction data.
 - Recipe: Example recipe strings were used, such as "`aggregate-stats :data_transfer_size :response_time total_size_mb total_time_sec`". Variations for average, median, p95, p99, and different output units were included where implemented.
 - Execution: The `TestingRig.execute(recipe, rows)` method was used to execute the recipe.
 - Expected Output: Assertions were used to verify that the output contains a single row with the correctly calculated aggregate values.

- Size Calculation: The `data_transfer_size` values were summed, converted to bytes, and then converted to Megabytes (MB) for the `total_size_mb` column. Consistency was maintained in the byte to MB conversion (either $1024*1024$ or $1000*1000$).
- Time Calculation: The `response_time` values were summed, converted to nanoseconds or milliseconds, and then converted to seconds for the `total_time_sec` column. If the average was implemented, the sum was divided by the number of rows before unit conversion.

4. AI Tools Usage

- DeepSeek - To break the assignment pdf into meaningful and optimized prompts.
- ChatGPT - For syntax and understanding the needs of the assignment
- Cursor - For assisted coding and debugging
- Gemini - Research for some of the bugs that i couldnt resolve

5. Deliverables

- The assignment is committed to GitHub.
- Modified `Directives.g4` file.
- New and modified Java source files within `wrangler-api` and `wrangler-core` modules.
- New and modified unit test files within `wrangler-core`.
- Evidence of successful build and test execution.
- `prompts.txt` file (if AI tooling was used).

6. Evaluation Criteria

- **6.1 Correctness:**
 - The new tokens are parsed correctly.
 - The directive computes aggregates accurately.
 - All tests pass.
- **6.2 Code Quality:**
 - The code is well-formatted, commented, and easy to understand.
 - Existing patterns are followed.
- **6.3 Robustness:**
 - The code handles edge cases, such as zero values, large numbers, and different unit cases.
- **6.4 Test Coverage:**
 - The new features are adequately tested.

Proof Of Work

```
[INFO] Recompiling the module because of added or removed source files.
[INFO] Compiling 1 source file with javac [debug target 1.8] to target/test-classes
[INFO]
[INFO] --- checkstyle:2.17:check (validate) @ wrangler-transform ---
[WARNING] Parameter 'sourceDirectory' is deprecated: instead use {@link #sourceDirectories}
[INFO]
[INFO] --- surefire:2.14.1:test (default-test) @ wrangler-transform ---
[INFO] -----
[INFO] Reactor Summary for Wrangler 4.12.0-SNAPSHOT:
[INFO]
[INFO] Wrangler ..... SUCCESS [ 28.111 s]
[INFO] Wrangler API ..... SUCCESS [ 7.602 s]
[INFO] Wrangler REST Protocol Classes ..... SUCCESS [ 12.088 s]
[INFO] Wrangler Core ..... SUCCESS [ 45.204 s]
[INFO] Wrangler Storage ..... SUCCESS [ 6.336 s]
[INFO] Wrangler Service ..... SUCCESS [ 12.889 s]
[INFO] Wrangler Testing Framework ..... SUCCESS [ 2.216 s]
[INFO] Wrangler Transform ..... SUCCESS [ 14.402 s]
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 02:11 min
[INFO] Finished at: 2025-04-13T12:17:23Z
[INFO] -----
```

The Modules that I was supposed to work on all **compile successfully** and **pass** all the test cases.